



RISC-V Vector C Intrinsic Specification Document

RISC-V Vector C Intrinsic Task Group

Version v0.13.0-draft, 2024-11-19: Draft

Table of Contents

Preamble	1
Copyright and license information	2
Contributors	3
1. Introduction	4
2. Test macro	5
3. Header file inclusion	6
4. Availability	7
5. Control of the vector extension programming model	8
5.1. Control of effective element width (EEW) and effective LMUL (EMUL)	8
5.2. Control of number of elements to be processed	8
5.3. Control of vector masking	9
5.4. Control of behavior of destination tail elements and destination inactive masked-off elements	9
5.5. Control of fixed-point rounding mode	9
5.6. Control of floating-point rounding mode	10
5.6.1. Implicit FP rounding mode intrinsics	10
5.6.2. Explicit FP rounding mode intrinsics	10
6. Naming scheme	12
6.1. Policy and masked naming scheme	12
6.2. Explicit (Non-overloaded) naming scheme	13
6.3. Exceptions in the explicit (non-overloaded) naming scheme	14
6.3.1. Scalar move instructions	14
6.3.2. Reduction instructions	14
6.3.3. Add-with-carry / Subtract-with-borrow instructions	14
6.3.4. vreinterpret , vlmul_trunc/vlmul_ext , and vset/vget	14
6.4. Implicit (Overloaded) naming scheme	15
6.5. Exceptions in the implicit (overloaded) naming scheme	16
6.5.1. Widening instructions	16
6.5.2. Type-convert instructions	16
6.5.3. vreinterpret , LMUL truncate/extension, and vset/vget	16
6.6. Un-supported intrinsics for implicit (overloaded) naming scheme	17
7. Type system	18
7.1. Integer types	18
7.2. Floating-point types	18
7.3. Mask types	19
7.4. Tuple type	19
8. Pseudo intrinsics	21
8.1. vsetvl	21
8.2. vsetvlmax	21
8.3. vreinterpret	21

8.4. <code>vundefined</code>	21
8.5. <code>vget</code>	21
8.6. <code>vset</code>	22
8.7. <code>vlmul_trunc</code>	22
8.8. <code>vlmul_ext</code>	22
8.9. <code>vcreate</code>	22
8.10. <code>vlenb</code>	22
9. Programming Notes	24
9.1. Agnostic value in the RVV C intrinsics	24
9.2. Copying vector register group contents	24
9.3. The passthrough (<code>vd</code>) argument in the intrinsics	24
9.4. Assumption of <code>vstart=0</code> for intrinsics users.	24
9.5. Assembly generated from the intrinsics	25
9.6. Bookkeeping of configurations	25
9.7. Strided load/store with stride of 0	25
9.8. Leveraging instructions with operand mnemonics of <code>vi</code>	25
9.9. Mixing inline assembly and intrinsics	25
9.10. The <code>new_vl</code> argument in fault-only-first load intrinsics.	25
9.11. Naming scheme	26
9.12. Control of the vector extension programming model	26
9.13. Type system	26
9.14. Psuedo intrinsics	26
10. References	27
11. Examples	29
11.1. Memory copy	29
11.2. SAXPY	29
11.3. Matrix multiplication	30
11.4. String copy	31
11.5. Control flow	31
11.6. Reduction and counting	32
Appendix A: Explicit (Non-overloaded) intrinsics	34
A.1. Vector Loads and Stores Intrinsics	34
A.1.1. Vector Unit-Stride Load Intrinsics	34
A.1.2. Vector Unit-Stride Store Intrinsics	36
A.1.3. Vector Mask Load/Store Intrinsics	39
A.1.4. Vector Strided Load Intrinsics	39
A.1.5. Vector Strided Store Intrinsics	42
A.1.6. Vector Indexed Load Intrinsics	46
A.1.7. Vector Indexed Store Intrinsics	71
A.1.8. Unit-stride Fault-Only-First Loads Intrinsics	96
A.2. Vector Loads and Stores Segment Intrinsics	100
A.2.1. Vector Unit-Stride Segment Load Intrinsics	100
A.2.2. Vector Unit-Stride Segment Store Intrinsics	124

A.2.3. Vector Strided Segment Load Intrinsics	135
A.2.4. Vector Strided Segment Store Intrinsics	149
A.2.5. Vector Indexed Segment Load Intrinsics	162
A.2.6. Vector Indexed Segment Store Intrinsics	279
A.3. Vector Integer Arithmetic Intrinsics	407
A.3.1. Vector Single-Width Integer Add and Subtract Intrinsics	407
A.3.2. Vector Widening Integer Add/Subtract Intrinsics	418
A.3.3. Vector Integer Widening Intrinsics	430
A.3.4. Vector Integer Extension Intrinsics	431
A.3.5. Vector Integer Add-with-Carry / Subtract-with-Borrow Intrinsics	433
A.3.6. Vector Bitwise Binary Logical Intrinsics	447
A.3.7. Vector Bitwise Unary Logical Intrinsics	459
A.3.8. Vector Single-Width Bit Shift Intrinsics	460
A.3.9. Vector Narrowing Integer Right Shift Intrinsics	468
A.3.10. Vector Integer Narrowing Intrinsics	471
A.3.11. Vector Integer Compare Intrinsics	472
A.3.12. Vector Integer Min/Max Intrinsics	498
A.3.13. Vector Single-Width Integer Multiply Intrinsics	506
A.3.14. Vector Integer Divide Intrinsics	516
A.3.15. Vector Widening Integer Multiply Intrinsics	524
A.3.16. Vector Single-Width Integer Multiply-Add Intrinsics	529
A.3.17. Vector Widening Integer Multiply-Add Intrinsics	551
A.3.18. Vector Integer Merge Intrinsics	559
A.3.19. Vector Integer Move Intrinsics	561
A.4. Vector Fixed-Point Arithmetic Intrinsics	563
A.4.1. Vector Single-Width Saturating Add and Subtract Intrinsics	563
A.4.2. Vector Single-Width Averaging Add and Subtract Intrinsics	571
A.4.3. Vector Single-Width Fractional Multiply with Rounding and Saturation Intrinsics	584
A.4.4. Vector Single-Width Scaling Shift Intrinsics	587
A.4.5. Vector Narrowing Fixed-Point Clip Intrinsics	592
A.5. Vector Floating-Point Intrinsics	597
A.5.1. Vector Single-Width Floating-Point Add/Subtract Intrinsics	597
A.5.2. Vector Widening Floating-Point Add/Subtract Intrinsics	607
A.5.3. Vector Single-Width Floating-Point Multiply/Divide Intrinsics	616
A.5.4. Vector Widening Floating-Point Multiply Intrinsics	626
A.5.5. Vector Single-Width Floating-Point Fused Multiply-Add Intrinsics	628
A.5.6. Vector Widening Floating-Point Fused Multiply-Add Intrinsics	667
A.5.7. Vector Floating-Point Square-Root Intrinsics	679
A.5.8. Vector Floating-Point Reciprocal Square-Root Estimate Intrinsics	681
A.5.9. Vector Floating-Point MIN/MAX Intrinsics	683
A.5.10. Vector Floating-Point Sign-Injection Intrinsics	687
A.5.11. Vector Floating-Point Absolute Value Intrinsics	692
A.5.12. Vector Floating-Point Compare Intrinsics	692

A.5.13. Vector Floating-Point Classify Intrinsics	702
A.5.14. Vector Floating-Point Merge Intrinsics	702
A.5.15. Vector Floating-Point Move Intrinsics	703
A.5.16. Single-Width Floating-Point/Integer Type-Convert Intrinsics	704
A.5.17. Widening Floating-Point/Integer Type-Convert Intrinsics	712
A.5.18. Narrowing Floating-Point/Integer Type-Convert Intrinsics	716
A.6. Vector Reduction Operations	724
A.6.1. Vector Single-Width Integer Reduction Intrinsics	724
A.6.2. Vector Widening Integer Reduction Intrinsics	740
A.6.3. Vector Single-Width Floating-Point Reduction Intrinsics	742
A.6.4. Vector Widening Floating-Point Reduction Intrinsics	749
A.7. Vector Mask Intrinsics	753
A.7.1. Vector Mask-Register Logical	753
A.7.2. Vector count population in mask vcpop.m	754
A.7.3. vfist find-first-set mask bit	755
A.7.4. vmsbf.m set-before-first mask bit	755
A.7.5. vmsif.m set-including-first mask bit	755
A.7.6. vmsof.m set-only-first mask bit	756
A.7.7. Vector Iota Intrinsics	756
A.7.8. Vector Element Index Intrinsics	757
A.8. Vector Permutation Intrinsics	757
A.8.1. Integer and Floating-Point Scalar Move Intrinsics	757
A.8.2. Vector Slideup Intrinsics	759
A.8.3. Vector Slidedown Intrinsics	763
A.8.4. Vector Slide1up and Slide1down Intrinsics	766
A.8.5. Vector Register Gather Intrinsics	773
A.8.6. Vector Compress Intrinsics	782
A.9. Miscellaneous Vector Utility Intrinsics	784
A.9.1. Get vl with specific vtype	784
A.9.2. Get VLMAX with specific vtype	784
A.9.3. Reinterpret Cast Conversion Intrinsics	785
A.9.4. Vector LMUL Extension Intrinsics	789
A.9.5. Vector LMUL Truncation Intrinsics	791
A.9.6. Vector Initialization Intrinsics	793
A.9.7. Vector Insertion Intrinsics	798
A.9.8. Vector Extraction Intrinsics	807
A.9.9. Vector Creation Intrinsics	811
Appendix B: Explicit (Non-overloaded) intrinsics, policy variants	822
B.1. Vector Loads and Stores Intrinsics	822
B.1.1. Vector Unit-Stride Load Intrinsics	822
B.1.2. Vector Unit-Stride Store Intrinsics	829
B.1.3. Vector Mask Load/Store Intrinsics	829
B.1.4. Vector Strided Load Intrinsics	829

B.1.5. Vector Strided Store Intrinsics	839
B.1.6. Vector Indexed Load Intrinsics	839
B.1.7. Vector Indexed Store Intrinsics	909
B.1.8. Unit-stride Fault-Only-First Loads Intrinsics	909
B.2. Vector Loads and Stores Segment Intrinsics	918
B.2.1. Vector Unit-Stride Segment Load Intrinsics	918
B.2.2. Vector Unit-Stride Segment Store Intrinsics	990
B.2.3. Vector Strided Segment Load Intrinsics	991
B.2.4. Vector Strided Segment Store Intrinsics	1032
B.2.5. Vector Indexed Segment Load Intrinsics	1033
B.2.6. Vector Indexed Segment Store Intrinsics	1389
B.3. Vector Integer Arithmetic Intrinsics	1389
B.3.1. Vector Single-Width Integer Add and Subtract Intrinsics	1389
B.3.2. Vector Widening Integer Add/Subtract Intrinsics	1421
B.3.3. Vector Integer Widening Intrinsics	1457
B.3.4. Vector Integer Extension Intrinsics	1461
B.3.5. Vector Integer Add-with-Carry / Subtract-with-Borrow Intrinsics	1467
B.3.6. Vector Integer Carry-out / Borrow-out Intrinsics	1473
B.3.7. Vector Bitwise Binary Logical Intrinsics	1473
B.3.8. Vector Bitwise Unary Logical Intrinsics	1509
B.3.9. Vector Single-Width Bit Shift Intrinsics	1513
B.3.10. Vector Narrowing Integer Right Shift Intrinsics	1537
B.3.11. Vector Integer Narrowing Intrinsics	1546
B.3.12. Vector Integer Compare Intrinsics	1549
B.3.13. Vector Integer Min/Max Intrinsics	1569
B.3.14. Vector Single-Width Integer Multiply Intrinsics	1594
B.3.15. Vector Integer Divide Intrinsics	1626
B.3.16. Vector Widening Integer Multiply Intrinsics	1651
B.3.17. Vector Single-Width Integer Multiply-Add Intrinsics	1664
B.3.18. Vector Widening Integer Multiply-Add Intrinsics	1716
B.3.19. Vector Integer Merge Intrinsics	1732
B.3.20. Vector Integer Move Intrinsics	1735
B.4. Vector Fixed-Point Arithmetic Intrinsics	1737
B.4.1. Vector Single-Width Saturating Add and Subtract Intrinsics	1737
B.4.2. Vector Single-Width Averaging Add and Subtract Intrinsics	1763
B.4.3. Vector Single-Width Fractional Multiply with Rounding and Saturation Intrinsics	1793
B.4.4. Vector Single-Width Scaling Shift Intrinsics	1800
B.4.5. Vector Narrowing Fixed-Point Clip Intrinsics	1815
B.5. Vector Floating-Point Intrinsics	1826
B.5.1. Vector Single-Width Floating-Point Add/Subtract Intrinsics	1826
B.5.2. Vector Widening Floating-Point Add/Subtract Intrinsics	1853
B.5.3. Vector Single-Width Floating-Point Multiply/Divide Intrinsics	1879
B.5.4. Vector Widening Floating-Point Multiply Intrinsics	1904

B.5.5. Vector Single-Width Floating-Point Fused Multiply-Add Intrinsics	1911
B.5.6. Vector Widening Floating-Point Fused Multiply-Add Intrinsics	1995
B.5.7. Vector Floating-Point Square-Root Intrinsics	2021
B.5.8. Vector Floating-Point Reciprocal Square-Root Estimate Intrinsics	2026
B.5.9. Vector Floating-Point MIN/MAX Intrinsics	2032
B.5.10. Vector Floating-Point Sign-Injection Intrinsics	2042
B.5.11. Vector Floating-Point Absolute Value Intrinsics	2056
B.5.12. Vector Floating-Point Compare Intrinsics	2058
B.5.13. Vector Floating-Point Classify Intrinsics	2066
B.5.14. Vector Floating-Point Merge Intrinsics	2068
B.5.15. Vector Floating-Point Move Intrinsics	2069
B.5.16. Single-Width Floating-Point/Integer Type-Convert Intrinsics	2070
B.5.17. Widening Floating-Point/Integer Type-Convert Intrinsics	2091
B.5.18. Narrowing Floating-Point/Integer Type-Convert Intrinsics	2103
B.6. Vector Reduction Operations	2125
B.6.1. Vector Single-Width Integer Reduction Intrinsics	2125
B.6.2. Vector Widening Integer Reduction Intrinsics	2145
B.6.3. Vector Single-Width Floating-Point Reduction Intrinsics	2147
B.6.4. Vector Widening Floating-Point Reduction Intrinsics	2157
B.7. Vector Mask Intrinsics	2162
B.7.1. Vector Mask-Register Logical	2162
B.7.2. Vector count population in mask vcpop.m	2162
B.7.3. vfist find-first-set mask bit	2162
B.7.4. vmsbf.m set-before-first mask bit	2163
B.7.5. vmsif.m set-including-first mask bit	2163
B.7.6. vmsof.m set-only-first mask bit	2163
B.7.7. Vector Iota Intrinsics	2163
B.7.8. Vector Element Index Intrinsics	2166
B.8. Vector Permutation Intrinsics	2167
B.8.1. Integer and Floating-Point Scalar Move Intrinsics	2167
B.8.2. Vector Slideup Intrinsics	2169
B.8.3. Vector Slidedown Intrinsics	2178
B.8.4. Vector Slide1up and Slide1down Intrinsics	2187
B.8.5. Vector Register Gather Intrinsics	2206
B.8.6. Vector Compress Intrinsics	2234
B.9. Miscellaneous Vector Utility Intrinsics	2236
B.9.1. Get vl with specific vtype	2236
B.9.2. Get VLMAX with specific vtype	2236
B.9.3. Reinterpret Cast Conversion Intrinsics	2237
B.9.4. Vector LMUL Extension Intrinsics	2237
B.9.5. Vector LMUL Truncation Intrinsics	2237
B.9.6. Vector Initialization Intrinsics	2237
B.9.7. Vector Insertion Intrinsics	2237

B.9.8. Vector Extraction Intrinsics	2237
B.9.9. Vector Creation Intrinsics	2237
Appendix C: Implicit (Overloaded) intrinsics	2238
C.1. Vector Loads and Stores Intrinsics	2238
C.1.1. Vector Unit-Stride Load Intrinsics	2238
C.1.2. Vector Unit-Stride Store Intrinsics	2239
C.1.3. Vector Mask Load/Store Intrinsics	2241
C.1.4. Vector Strided Load Intrinsics	2241
C.1.5. Vector Strided Store Intrinsics	2243
C.1.6. Vector Indexed Load Intrinsics	2245
C.1.7. Vector Indexed Store Intrinsics	2264
C.1.8. Unit-stride Fault-Only-First Loads Intrinsics	2287
C.2. Vector Loads and Stores Segment Intrinsics	2289
C.2.1. Vector Unit-Stride Segment Load Intrinsics	2289
C.2.2. Vector Unit-Stride Segment Store Intrinsics	2299
C.2.3. Vector Strided Segment Load Intrinsics	2307
C.2.4. Vector Strided Segment Store Intrinsics	2314
C.2.5. Vector Indexed Segment Load Intrinsics	2327
C.2.6. Vector Indexed Segment Store Intrinsics	2432
C.3. Vector Integer Arithmetic Intrinsics	2538
C.3.1. Vector Single-Width Integer Add and Subtract Intrinsics	2538
C.3.2. Vector Widening Integer Add/Subtract Intrinsics	2547
C.3.3. Vector Integer Widening Intrinsics	2558
C.3.4. Vector Integer Extension Intrinsics	2559
C.3.5. Vector Integer Add-with-Carry / Subtract-with-Borrow Intrinsics	2560
C.3.6. Vector Bitwise Binary Logical Intrinsics	2570
C.3.7. Vector Bitwise Unary Logical Intrinsics	2580
C.3.8. Vector Single-Width Bit Shift Intrinsics	2581
C.3.9. Vector Narrowing Integer Right Shift Intrinsics	2588
C.3.10. Vector Integer Narrowing Intrinsics	2590
C.3.11. Vector Integer Compare Intrinsics	2591
C.3.12. Vector Integer Min/Max Intrinsics	2610
C.3.13. Vector Single-Width Integer Multiply Intrinsics	2617
C.3.14. Vector Integer Divide Intrinsics	2626
C.3.15. Vector Widening Integer Multiply Intrinsics	2633
C.3.16. Vector Single-Width Integer Multiply-Add Intrinsics	2637
C.3.17. Vector Widening Integer Multiply-Add Intrinsics	2657
C.3.18. Vector Integer Merge Intrinsics	2663
C.3.19. Vector Integer Move Intrinsics	2665
C.4. Vector Fixed-Point Arithmetic Intrinsics	2666
C.4.1. Vector Single-Width Saturating Add and Subtract Intrinsics	2666
C.4.2. Vector Single-Width Averaging Add and Subtract Intrinsics	2673
C.4.3. Vector Single-Width Fractional Multiply with Rounding and Saturation Intrinsics	2684

C.4.4. Vector Single-Width Scaling Shift Intrinsics	2687
C.4.5. Vector Narrowing Fixed-Point Clip Intrinsics	2692
C.5. Vector Floating-Point Intrinsics	2696
C.5.1. Vector Single-Width Floating-Point Add/Subtract Intrinsics	2696
C.5.2. Vector Widening Floating-Point Add/Subtract Intrinsics	2704
C.5.3. Vector Single-Width Floating-Point Multiply/Divide Intrinsics	2712
C.5.4. Vector Widening Floating-Point Multiply Intrinsics	2720
C.5.5. Vector Single-Width Floating-Point Fused Multiply-Add Intrinsics	2722
C.5.6. Vector Widening Floating-Point Fused Multiply-Add Intrinsics	2750
C.5.7. Vector Floating-Point Square-Root Intrinsics	2759
C.5.8. Vector Floating-Point Reciprocal Square-Root Estimate Intrinsics	2760
C.5.9. Vector Floating-Point MIN/MAX Intrinsics	2762
C.5.10. Vector Floating-Point Sign-Injection Intrinsics	2765
C.5.11. Vector Floating-Point Absolute Value Intrinsics	2769
C.5.12. Vector Floating-Point Compare Intrinsics	2769
C.5.13. Vector Floating-Point Classify Intrinsics	2777
C.5.14. Vector Floating-Point Merge Intrinsics	2777
C.5.15. Vector Floating-Point Move Intrinsics	2778
C.5.16. Single-Width Floating-Point/Integer Type-Convert Intrinsics	2778
C.5.17. Widening Floating-Point/Integer Type-Convert Intrinsics	2784
C.5.18. Narrowing Floating-Point/Integer Type-Convert Intrinsics	2787
C.6. Vector Reduction Operations	2793
C.6.1. Vector Single-Width Integer Reduction Intrinsics	2793
C.6.2. Vector Widening Integer Reduction Intrinsics	2804
C.6.3. Vector Single-Width Floating-Point Reduction Intrinsics	2806
C.6.4. Vector Widening Floating-Point Reduction Intrinsics	2810
C.7. Vector Mask Intrinsics	2813
C.7.1. Vector Mask-Register Logical	2813
C.7.2. Vector count population in mask vcpop.m	2814
C.7.3. vfist find-first-set mask bit	2814
C.7.4. vmsbf.m set-before-first mask bit	2815
C.7.5. vmsif.m set-including-first mask bit	2815
C.7.6. vmsof.m set-only-first mask bit	2815
C.7.7. Vector Iota Intrinsics	2816
C.7.8. Vector Element Index Intrinsics	2816
C.8. Vector Permutation Intrinsics	2816
C.8.1. Integer and Floating-Point Scalar Move Intrinsics	2816
C.8.2. Vector Slideup Intrinsics	2817
C.8.3. Vector Slidedown Intrinsics	2820
C.8.4. Vector Slide1up and Slide1down Intrinsics	2823
C.8.5. Vector Register Gather Intrinsics	2828
C.8.6. Vector Compress Intrinsics	2836
C.9. Miscellaneous Vector Utility Intrinsics	2837

C.9.1. Get v1 with specific vtype	2837
C.9.2. Get VLMAX with specific vtype	2837
C.9.3. Reinterpret Cast Conversion Intrinsics	2837
C.9.4. Vector LMUL Extension Intrinsics	2842
C.9.5. Vector LMUL Truncation Intrinsics	2844
C.9.6. Vector Initialization Intrinsics	2846
C.9.7. Vector Insertion Intrinsics	2846
C.9.8. Vector Extraction Intrinsics	2852
C.9.9. Vector Creation Intrinsics	2856
Appendix D: Implicit (Overloaded) intrinsics, policy variants	2857
D.1. Vector Loads and Stores Intrinsics	2857
D.1.1. Vector Unit-Stride Load Intrinsics	2857
D.1.2. Vector Unit-Stride Store Intrinsics	2863
D.1.3. Vector Mask Load/Store Intrinsics	2863
D.1.4. Vector Strided Load Intrinsics	2863
D.1.5. Vector Strided Store Intrinsics	2870
D.1.6. Vector Indexed Load Intrinsics	2870
D.1.7. Vector Indexed Store Intrinsics	2931
D.1.8. Unit-stride Fault-Only-First Loads Intrinsics	2931
D.2. Vector Loads and Stores Segment Intrinsics	2939
D.2.1. Vector Unit-Stride Segment Load Intrinsics	2939
D.2.2. Vector Unit-Stride Segment Store Intrinsics	3002
D.2.3. Vector Strided Segment Load Intrinsics	3002
D.2.4. Vector Strided Segment Store Intrinsics	3038
D.2.5. Vector Indexed Segment Load Intrinsics	3038
D.2.6. Vector Indexed Segment Store Intrinsics	3331
D.3. Vector Integer Arithmetic Intrinsics	3331
D.3.1. Vector Single-Width Integer Add and Subtract Intrinsics	3331
D.3.2. Vector Widening Integer Add/Subtract Intrinsics	3360
D.3.3. Vector Integer Widening Intrinsics	3389
D.3.4. Vector Integer Extension Intrinsics	3392
D.3.5. Vector Integer Add-with-Carry / Subtract-with-Borrow Intrinsics	3398
D.3.6. Vector Integer Carry-out / Borrow-out Intrinsics	3403
D.3.7. Vector Bitwise Binary Logical Intrinsics	3404
D.3.8. Vector Bitwise Unary Logical Intrinsics	3435
D.3.9. Vector Single-Width Bit Shift Intrinsics	3439
D.3.10. Vector Narrowing Integer Right Shift Intrinsics	3460
D.3.11. Vector Integer Narrowing Intrinsics	3468
D.3.12. Vector Integer Compare Intrinsics	3471
D.3.13. Vector Integer Min/Max Intrinsics	3487
D.3.14. Vector Single-Width Integer Multiply Intrinsics	3508
D.3.15. Vector Integer Divide Intrinsics	3534
D.3.16. Vector Widening Integer Multiply Intrinsics	3555

D.3.17. Vector Single-Width Integer Multiply-Add Intrinsics	3566
D.3.18. Vector Widening Integer Multiply-Add Intrinsics	3608
D.3.19. Vector Integer Merge Intrinsics	3621
D.3.20. Vector Integer Move Intrinsics	3623
D.4. Vector Fixed-Point Arithmetic Intrinsics	3625
D.4.1. Vector Single-Width Saturating Add and Subtract Intrinsics	3625
D.4.2. Vector Single-Width Averaging Add and Subtract Intrinsics	3646
D.4.3. Vector Single-Width Fractional Multiply with Rounding and Saturation Intrinsics	3667
D.4.4. Vector Single-Width Scaling Shift Intrinsics	3672
D.4.5. Vector Narrowing Fixed-Point Clip Intrinsics	3683
D.5. Vector Floating-Point Intrinsics	3691
D.5.1. Vector Single-Width Floating-Point Add/Subtract Intrinsics	3691
D.5.2. Vector Widening Floating-Point Add/Subtract Intrinsics	3711
D.5.3. Vector Single-Width Floating-Point Multiply/Divide Intrinsics	3734
D.5.4. Vector Widening Floating-Point Multiply Intrinsics	3753
D.5.5. Vector Single-Width Floating-Point Fused Multiply-Add Intrinsics	3758
D.5.6. Vector Widening Floating-Point Fused Multiply-Add Intrinsics	3822
D.5.7. Vector Floating-Point Square-Root Intrinsics	3843
D.5.8. Vector Floating-Point Reciprocal Square-Root Estimate Intrinsics	3846
D.5.9. Vector Floating-Point MIN/MAX Intrinsics	3852
D.5.10. Vector Floating-Point Sign-Injection Intrinsics	3859
D.5.11. Vector Floating-Point Absolute Value Intrinsics	3871
D.5.12. Vector Floating-Point Compare Intrinsics	3872
D.5.13. Vector Floating-Point Classify Intrinsics	3878
D.5.14. Vector Floating-Point Merge Intrinsics	3879
D.5.15. Vector Floating-Point Move Intrinsics	3880
D.5.16. Single-Width Floating-Point/Integer Type-Convert Intrinsics	3881
D.5.17. Widening Floating-Point/Integer Type-Convert Intrinsics	3898
D.5.18. Narrowing Floating-Point/Integer Type-Convert Intrinsics	3909
D.6. Vector Reduction Operations	3927
D.6.1. Vector Single-Width Integer Reduction Intrinsics	3928
D.6.2. Vector Widening Integer Reduction Intrinsics	3943
D.6.3. Vector Single-Width Floating-Point Reduction Intrinsics	3946
D.6.4. Vector Widening Floating-Point Reduction Intrinsics	3953
D.7. Vector Mask Intrinsics	3957
D.7.1. Vector Mask-Register Logical	3957
D.7.2. Vector count population in mask vcpop.m	3957
D.7.3. vf first find-first-set mask bit	3957
D.7.4. vmsbf.m set-before-first mask bit	3957
D.7.5. vmsif.m set-including-first mask bit	3957
D.7.6. vmsof.m set-only-first mask bit	3957
D.7.7. Vector Iota Intrinsics	3958
D.7.8. Vector Element Index Intrinsics	3960

D.8. Vector Permutation Intrinsics	3962
D.8.1. Integer and Floating-Point Scalar Move Intrinsics	3962
D.8.2. Vector Slideup Intrinsics	3963
D.8.3. Vector Slidedown Intrinsics	3970
D.8.4. Vector Slide1up and Slide1down Intrinsics	3977
D.8.5. Vector Register Gather Intrinsics	3991
D.8.6. Vector Compress Intrinsics	4015
D.9. Miscellaneous Vector Utility Intrinsics	4016
D.9.1. Get vl with specific vtype	4016
D.9.2. Get VLMAX with specific vtype	4017
D.9.3. Reinterpret Cast Conversion Intrinsics	4017
D.9.4. Vector LMUL Extension Intrinsics	4017
D.9.5. Vector LMUL Truncation Intrinsics	4017
D.9.6. Vector Initialization Intrinsics	4017
D.9.7. Vector Insertion Intrinsics	4017
D.9.8. Vector Extraction Intrinsics	4017
D.9.9. Vector Creation Intrinsics	4017
Bibliography	4018

Preamble



This document is in the [Development state](#)

Expect potential changes. This draft specification is likely to evolve before it is accepted as a standard. Implementations based on this draft may not conform to the future standard.

Copyright and license information

This specification is licensed under the Creative Commons Attribution 4.0 International License (CC-BY 4.0). The full license text is available at creativecommons.org/licenses/by/4.0/.

Copyright 2023 by RISC-V International.

Contributors

This RISC-V specification has been contributed to directly or indirectly by (in alphabetical order):

Contributors to all versions of the spec in alphabetical order:

Brandon Wu, Craig Topper, Eop Chen, HanKuan Chen, HsiangKai Wang, Jerry Zhang Jian, Kito Cheng, Nick Knight, Olaf Bernstein, Roger Ferrer Ibanez, Yi-Hsiu Hsu, Zakk Chen

Chapter 1. Introduction

The RISC-V vector C intrinsics provide users interfaces in the C language level to directly leverage the RISC-V "V" extension ([RISC-V "V" Vector Extension, n.d.](#)) (also abbreviated as "RVV"), with assistance from the compiler in handling instruction scheduling and register allocation. The intrinsics also aim to free users from responsibility of maintaining the correct configuration settings for the vector instruction executions.

This document uses the term "RVV" as an abbreviation for the RISC-V "V" extension. This document uses the term "the RVV specification" to indicate the RISC-V "V" extension specification.

Chapter 2. Test macro

The `__riscv_v_intrinsic` macro is the C macro to test the compiler's support for the RISC-V "V" extension intrinsics.

This macro should be defined even if the vector extension is not enabled.

The value of the test macro is defined as its version, which is computed using the following formula. The formula is identical to what is defined in the RISC-V C API specification ([RISC-V C API Specification, n.d.](#)).

```
<MAJOR_VERSION> * 1,000,000 + <MINOR_VERSION> * 1,000 + <REVISION_VERSION>
```

For example, the v1.0 version should define the macro with value **1000000**.

Chapter 3. Header file inclusion

To leverage the intrinsics in the toolchain, the header `<riscv_vector.h>` needs to be included. We suggest guarding the inclusion with the test macro.

```
#if __riscv_v_intrinsic >= 1000000
#include <riscv_vector.h>
#endif /* __riscv_v_intrinsic */
```

Chapter 4. Availability

With `<riscv_vector.h>` included, availability of intrinsic variants depends on the required architecture of their corresponding vector instructions. The supported architecture is specified to the compiler using the `-march` option ([User Guide for RISC-V Target, n.d.](#); [RISC-V Options, n.d.](#)).

The standard vector extensions ([RISC-V "V" Vector Extension, n.d.](#)) provides a set of smaller extensions for embedded use. Please check out the **Zve** extensions ([RISC-V "V" Vector Extension, n.d.](#)) for the varying degree of support.

For example, RVV type `vint64m1_t` and `__riscv_vle64_v_i64m1` are not available under architecture `rv64gc_zve32x`.

Chapter 5. Control of the vector extension programming model

The intrinsics allow users to control the fields in **vtype**, as well as the rounding modes for fixed-point (**vxrm**) and floating-point (**frm**) vector computations.

In this chapter, we cover how the intrinsics embed the control of **vtype** fields in the function names. Please see [Chapter 6](#) for the rules described in this chapter.

5.1. Control of effective element width (EEW) and effective LMUL (EMUL)

The RISC-V vector intrinsics' data types are strongly-typed. The vector intrinsics encode the EEW (effective-element-width) and EMUL (effective LMUL) of the destination vector register in the suffix of the function name. Users can expect the results of the vector instruction intrinsics are computed under the specified EEW and EMUL.

To see the full list of data types for the intrinsics, please see [Chapter 7](#).

In the following example, the intrinsic will produce the semantic of a **vadd.vv** instruction, with source vector operands of **EEW=32** and **EMUL=1**, which can be observed through provided RVV data type; and produce results for the destination vector operand with **EEW=32** and **EMUL=1**.

```
vint32m1_t __riscv_vadd_vv_i32m1(vint32m1_t vs2, vint32m1_t vs1, size_t vl);
```

In the following example, the intrinsic will produce the semantic of a **vadd.vv** instruction, with source vector operands of **EEW=16** and **EMUL=1/2**, which can be observed through provided RVV data type; and produce results for the destination vector operand with **EEW=32** and **EMUL=1**.

```
vint32m1_t __riscv_vwadd_vv_i32m1(vint16mf2_t vs2, vint16mf2_t vs1, size_t vl);
```

5.2. Control of number of elements to be processed

The intrinsics do not directly expose the vector length control register to the intrinsics programmer. The intrinsics programmer specifies an "application vector length (AVL)" using the argument **size_t vl**. The implementation is responsible to set the correct value into the underlying vector length control register (**vl**) given the informed AVL.



*The intrinsics for instructions that behave the same with different **vl** settings (e.g. **vmv.s.x**) do not have a **size_t vl** argument.*



*The actual value written to the **vl** control register is an implementation defined behavior and is typically not known until runtime. The actual setting of **vl**, given the provided AVL through the parameter, follows the rules in the RVV specification. The number of elements processed can be obtained through the **_riscv_vsetvl*** intrinsics [Section 8.1](#).*

5.3. Control of vector masking

Instructions that are available for masking have masked variant intrinsics.

The intrinsics fuse the control of vector masking (**vm**) together with the control for policy behavior (**vta**, **vma**) in the same suffix. Please check out [Section 5.4](#) and [Section 6.1](#) for the exact suffix that specifies a masked/unmasked vector operation along with its policy behavior.

5.4. Control of behavior of destination tail elements and destination inactive masked-off elements

The behavior of destination tail elements and destination inactive masked-off elements is controlled by the **vta** and **vma** bits.

Given the general assumption that target audience of the intrinsics are high performance cores, and an "undisturbed" policy will generally slow down an out-of-order core, the intrinsics have a default policy scheme of tail-agnostic and mask-agnostic (that is, **vta**=1 and **vma**=1).

The intrinsics fuse the control of vector masking (**vm**) together with the control for policy behavior (**vta**, **vma**) in the same suffix. Please checkout [Section 5.3](#) and [Section 6.1](#) for the exact suffix that specifies a masked/unmasked vector operation along with its policy behavior.

5.5. Control of fixed-point rounding mode

For the fixed-point intrinsics, representing the fixed-point arithmetic instructions, the **vxrm** argument of the intrinsics indicates the rounding mode (**vxrm**) control.

The **vxrm** argument is required to be a constant integer expression. The implementation should provide the following **enum** that maps to the defined rounding mode values under Table 4 of the RVV specification.

```
enum __RISC_VXRM {
    __RISC_VXRM_RNU = 0,
    __RISC_VXRM_RNE = 1,
    __RISC_VXRM_RDN = 2,
    __RISC_VXRM_ROD = 3,
};
```



*Rounding mode does not affect the computations of **vsadd**, **vsaddu**, **vssub**, and **vssubu**; therefore, the intrinsics for these instructions do not have the **vxrm** argument.*



*The RISC-V psABI ([RISC-V Calling Conventions: Vector, n.d.](#)) states that **vxrm** is not preserved across calls. Optimization for reducing the number of redundant writes to **vxrm** is a compiler and system specific issue.*



*This specification does not provide support for manipulating the **vxsat** CSR. Since **vxsat** is not needed by a large majority of fixed-point code, we believe this specification is broadly useful as-is. Nevertheless, we expect that a future extension will define an additional set of fixed-point intrinsics that update **vxsat** in a specified manner, along with intrinsics to*

explicitly read and write **vxsat**. These new intrinsics would be interoperable with the intrinsics in this specification.

The value of the **vxsat** after a fixed-point intrinsic is **UNSPECIFIED**.

5.6. Control of floating-point rounding mode

For the floating-point intrinsics, representing the floating-point arithmetic instructions, the intrinsics have two variants: *Implicit FP rounding mode* and *Explicit FP Rounding mode* intrinsics.



Control of the floating-point accrued exceptions flag fields (**fflag**) ([RISC-V "F" Vector Extension, n.d.](#)) is not yet covered in the vector intrinsics v1.0. We plan to support it in follow-up versions in a compatible way with existing intrinsics in v1.0.

5.6.1. Implicit FP rounding mode intrinsics

The implicit FP rounding mode intrinsics behave like any C-language floating-point expressions, using the default rounding mode when **FENV_ACCESS** is off, and using the **fenv** dynamic rounding mode when **FENV_ACCESS** is on.



Both GNU and LLVM compilers generate scalar floating-point instructions using dynamic rounding mode, relying on the environment initialization to set **frm** to **RNE** (specified as "roundTiesToEven" in IEEE-754 (a.k.a. IEC 60559)) ([ANSI/IEEE Std 754-2008, IEEE Standard for Floating-Point Arithmetic, 2008](#)).



The implicit FP rounding mode intrinsics are intended to be used regardless of **FENV_ACCESS**. They are provided when **FENV_ACCESS** is on for the (few) programmers who are already using **fenv**; and they are provided when **FENV_ACCESS** is off for the (vast majority of) programmers who prefer the default rounding mode.

5.6.2. Explicit FP rounding mode intrinsics

The explicit FP rounding mode intrinsics contain the **frm** argument which indicates the rounding mode (**frm**) ([RISC-V "F" Vector Extension, n.d.](#)) control. The floating-point intrinsics with the **frm** argument are followed by an **_rm** suffix in the function name.

The **frm** argument is required to be a constant integer expression. The implementation should provide the following **enum** that maps to the defined rounding mode values under RISC-V ISA Manual Table 8.1 ([RISC-V Calling Conventions: Vector, n.d.](#)).

```
enum __RISCV_FRM {
    __RISCV_FRM_RNE = 0,
    __RISCV_FRM_RTZ = 1,
    __RISCV_FRM_RDN = 2,
    __RISCV_FRM_RUP = 3,
    __RISCV_FRM_RMM = 4,
};
```



The explicit FP rounding mode intrinsics are intended to be used when **FENV_ACCESS** is off, enabling more aggressive optimization while still providing the programmer with control over the rounding mode. Using explicit FP rounding mode intrinsics when **FENV_ACCESS** is on will still work correctly, but is expected to lead to extra saving/restoring of **frm**, that

*could be avoided by using **fenv** functionality and implicit FP rounding mode intrinsics.*

Chapter 6. Naming scheme

The naming scheme of the intrinsics expresses the users' control of fields in **vtype**, **vl**, and rounding modes for fixed-point and floating-point vector computations. For details of these CSR controls, please see [\[control-of-vector-programming-mode\]](#).

As mentioned in [Section 5.3](#) and [Section 5.4](#), the intrinsics fuses the control of **vm**, **vta**, and **vma** into the same suffix. [Section 6.1](#) enumerates the exact suffixes. You may find where these suffixes are appended in [Section 6.2](#).

The intrinsics can be split into two major types, called "explicit (non-overloaded) intrinsics" and "implicit (overloaded) intrinsics".

The explicit (non-overloaded) intrinsics embed the control described in [Chapter 5](#) in the function name. This scheme gives intrinsic codebase more readability as the execution states are explicitly specified in the code.

The implicit (overloaded) intrinsics, on the contrary, omit the explicit specifications for **vtype** control. The implicit (overloaded) intrinsics aim to provide a generic interface to let users put values of different EEW and EMUL as the input argument.

This section covers the general naming rule of the two types of intrinsics accordingly. Then, this section also enumerates the exceptions and the rationales behind them in [Section 6.3](#) and [Section 6.5](#).

6.1. Policy and masked naming scheme

With the default policy scheme mentioned under [Section 5.4](#), each intrinsic provides corresponding variants for their available control of **vm**, **vta** and **vma**. The following list enumerates the possible suffixes.

- No suffix: Represents an unmasked (**vm=1**) vector operation with tail-agnostic (**vta=1**)
- **_tu** suffix: Represents an unmasked (**vm=1**) vector operation with tail-undisturbed (**vta=0**) policy
- **_m** suffix: Represents a masked (**vm=0**) vector operation with tail-agnostic (**vta=1**), mask-agnostic (**vma=1**) policy
- **_tum** suffix: Represents a masked (**vm=0**) vector operation with tail-undisturbed (**vta=0**), mask-agnostic (**vma=1**) policy
- **_mu** suffix: Represents a masked (**vm=0**) vector operation with tail-agnostic (**vta=1**), mask-undisturbed (**vma=0**) policy
- **_tumu** suffix: Represents a masked (**vm=0**) vector operation with tail-undisturbed (**vta=0**), mask-undisturbed (**vma=0**) policy

Using **vadd** with EEW=32 and EMUL=1 as an example, the variants are:

```
// vm=1, vta=1
vint32m1_t __riscv_vadd_vv_i32m1(vint32m1_t vs2, vint32m1_t vs1, size_t vl);
// vm=1, vta=0
vint32m1_t __riscv_vadd_vv_i32m1_tu(vint32m1_t vd, vint32m1_t vs2,
                                   vint32m1_t vs1, size_t vl);
// vm=0, vta=1, vma=1
vint32m1_t __riscv_vadd_vv_i32m1_m(vbool32_t vm, vint32m1_t vs2, vint32m1_t vs1,
                                   size_t vl);
```



```
// vm=0, vta=0, vma=1
vint32m1_t __riscv_vadd_vv_i32m1_tum(vbool32_t vm, vint32m1_t vd,
                                     vint32m1_t vs2, vint32m1_t vs1, size_t vl);

// vm=0, vta=1, vma=0
vint32m1_t __riscv_vadd_vv_i32m1_mu(vbool32_t vm, vint32m1_t vd, vint32m1_t vs2,
                                    vint32m1_t vs1, size_t vl);

// vm=0, vta=0, vma=0
vint32m1_t __riscv_vadd_vv_i32m1_tumu(vbool32_t vm, vint32m1_t vd,
                                       vint32m1_t vs2, vint32m1_t vs1,
                                       size_t vl);
```



When policy is set to "agnostic", there is no guarantee of what will be in the tail/masked-off elements. Under this policy, users should not assume the values within to be deterministic.



Pseudo intrinsics mentioned under [Chapter 8](#) do not map to real vector instructions. Therefore these intrinsics are not affected by the policy setting, nor do they have intrinsic variants of the suffixes listed above.

6.2. Explicit (Non-overloaded) naming scheme

In general, the intrinsics are encoded as the following. The intrinsics under this naming scheme are the "non-overloaded intrinsics", which in parallel we have the "overloaded intrinsics" defined under [Section 6.4](#).

The naming rules are as follows.

```
__riscv_{V_INSTRUCTION_MNEMONIC}_{OPERAND_MNEMONIC}_{RETURN_TYPE}_{ROUND_MODE}_{POLICY}{(...)}
```

- **OPERAND_MNEMONIC** are like **v**, **vv**, **vx**, **vs**, **vvm**, **vxm**
- **RETURN_TYPE** depends on whether the return type of the vector instruction is a mask register...
 - For intrinsics that represents instructions with a non-mask destination register:
 - **EEW** is one of **i8** | **i16** | **i32** | **i64** | **u8** | **u16** | **u32** | **u64** | **f16** | **f32** | **f64**.
 - **EMUL** is one of **m1** | **m2** | **m4** | **m8** | **mf2** | **mf4** | **mf8**.
 - [Chapter 7](#) explains the limited enumeration of EEW-EMUL pairs.
 - For intrinsics that represent intrinsics with a mask destination register:
 - **RETURN_TYPE** is one of **b1** | **b2** | **b4** | **b8** | **b16** | **b32** | **b64**, which is derived from the ratio EEW/EMUL.
- **V_INSTRUCTION_MNEMONIC** are like **vadd**, **vfmac**, **vsadd**.
- **ROUND_MODE** is the **_rm** suffix mentioned in [Section 5.6.2](#). Other intrinsics do not have this suffix.
- **POLICY** are enumerated under [Section 6.1](#).

The general naming scheme is not sufficient to express all intrinsics. The exceptions are enumerated in the proceeding section [Section 6.3](#).

6.3. Exceptions in the explicit (non-overloaded) naming scheme

This section enumerates the exceptions in the explicit (non-overloaded) naming scheme.

6.3.1. Scalar move instructions

Only encoding the return type will cause naming collisions for the permutation instruction intrinsics. The intrinsics encode the input vector type and the output scalar type in the suffix.

```
int8_t __riscv_vmv_x_s_i8m1_i8 (vint8m1_t vs2, size_t vl);
int8_t __riscv_vmv_x_s_i8m2_i8 (vint8m2_t vs2, size_t vl);
int8_t __riscv_vmv_x_s_i8m4_i8 (vint8m4_t vs2, size_t vl);
int8_t __riscv_vmv_x_s_i8m8_i8 (vint8m8_t vs2, size_t vl);
```

6.3.2. Reduction instructions

Only encoding the return type will cause naming collisions for the reduction instruction intrinsics. The intrinsics encode the input vector type and the output vector type in the suffix.

```
vint8m1_t __riscv_vredsum_vs_i8m1_i8m1(vint8m1_t vs2, vint8m1_t vs1, size_t vl);
vint8m1_t __riscv_vredsum_vs_i8m2_i8m1(vint8m2_t vs2, vint8m1_t vs1, size_t vl);
vint8m1_t __riscv_vredsum_vs_i8m4_i8m1(vint8m4_t vs2, vint8m1_t vs1, size_t vl);
vint8m1_t __riscv_vredsum_vs_i8m8_i8m1(vint8m8_t vs2, vint8m1_t vs1, size_t vl);
```

6.3.3. Add-with-carry / Subtract-with-borrow instructions

Only encoding the return type will cause naming collisions for the reduction instruction intrinsics. The intrinsics encode the input vector type and the output mask vector type in the suffix.

```
vbool64_t __riscv_vmacd_vvm_i8mf8_b64(vint8mf8_t vs2, vint8mf8_t vs1,
                                       vbool64_t v0, size_t vl);
vbool64_t __riscv_vmacd_vvm_i16mf4_b64(vint16mf4_t vs2, vint16mf4_t vs1,
                                       vbool64_t v0, size_t vl);
vbool64_t __riscv_vmacd_vvm_i32mf2_b64(vint32mf2_t vs2, vint32mf2_t vs1,
                                       vbool64_t v0, size_t vl);
vbool64_t __riscv_vmacd_vvm_i64m1_b64(vint64m1_t vs2, vint64m1_t vs1,
                                       vbool64_t v0, size_t vl);
```

6.3.4. vreinterpret, vlmul_trunc/vlmul_ext, and vset/vget

Only encoding the return type will cause naming collisions for these pseudo instructions. The intrinsics encode the input vector type before the return type in the suffix.

The following shows an example with `__riscv_vreinterpret_v` of `vint32m1_t` input vector type.

```
vfloat32m1_t __riscv_vreinterpret_v_i32m1_f32m1 (vint32m1_t src);
vuint32m1_t __riscv_vreinterpret_v_i32m1_u32m1 (vint32m1_t src);
vint8m1_t __riscv_vreinterpret_v_i32m1_i8m1 (vint32m1_t src);
vint16m1_t __riscv_vreinterpret_v_i32m1_i16m1 (vint32m1_t src);
vint64m1_t __riscv_vreinterpret_v_i32m1_i64m1 (vint32m1_t src);
vbool64_t __riscv_vreinterpret_v_i32m1_b64 (vint32m1_t src);
```

```
vbool32_t __riscv_vreinterpret_v_i32m1_b32 (vint32m1_t src);
vbool16_t __riscv_vreinterpret_v_i32m1_b16 (vint32m1_t src);
vbool8_t __riscv_vreinterpret_v_i32m1_b8 (vint32m1_t src);
vbool4_t __riscv_vreinterpret_v_i32m1_b4 (vint32m1_t src);
```

6.4. Implicit (Overloaded) naming scheme

The implicit (overloaded) interface aims to provide a generic interface that takes values of different EEW and EMUL as the input. Therefore, the implicit intrinsics omit the EEW and EMUL encoded in the function name. The `_rm` prefix for explicit FP rounding mode intrinsics ([Section 5.6](#)) is also omitted. The intrinsics under this scheme are the "overloaded intrinsics", which in parallel we have the "non-overloaded intrinsics" defined under [Section 6.2](#).

Take the vector addition (**vadd**) instruction intrinsics as an example, stripping off the operand mnemonics and encoded EEW, EMUL information, the intrinsics provides the following overloaded interfaces.

```
vint32m1_t __riscv_vadd(vint32m1_t v0, vint32m1_t v1, size_t vl);
vint16m4_t __riscv_vadd(vint16m4_t v0, vint16m4_t v1, size_t vl);
```

Since the main intent is to let the users put different value(s) of EEW and EMUL as input argument(s), the overloaded intrinsics do not omit the policy suffix. That is, the suffix listed under [Section 5.4](#) is not omitted and is still encoded in the function name.

The masked variants with the default policy shares the same interface with the unmasked variants with the default policy. They do not have any trailing suffixes.

Take the vector floating-point add (**vfadd**) as an example, the intrinsics provides the following overloaded interfaces.

```
vfloat32m1_t __riscv_vfadd(vfloat32m1_t vs2, vfloat32m1_t vs1, size_t vl);
vfloat32m1_t __riscv_vfadd(vbool32_t vm, vfloat32m1_t vs2, vfloat32m1_t vs1,
                          size_t vl);
vfloat32m1_t __riscv_vfadd(vfloat32m1_t vs2, vfloat32m1_t vs1, unsigned int frm,
                          size_t vl);
vfloat32m1_t __riscv_vfadd(vbool32_t vm, vfloat32m1_t vs2, vfloat32m1_t vs1,
                          unsigned int frm, size_t vl);
vfloat32m1_t __riscv_vfadd_tu(vfloat32m1_t vd, vfloat32m1_t vs2,
                             vfloat32m1_t vs1, size_t vl);
vfloat32m1_t __riscv_vfadd_tum(vbool32_t vm, vfloat32m1_t vd, vfloat32m1_t vs2,
                              vfloat32m1_t vs1, size_t vl);
vfloat32m1_t __riscv_vfadd_tumu(vbool32_t vm, vfloat32m1_t vd, vfloat32m1_t vs2,
                               vfloat32m1_t vs1, size_t vl);
vfloat32m1_t __riscv_vfadd_mu(vbool32_t vm, vfloat32m1_t vd, vfloat32m1_t vs2,
                              vfloat32m1_t vs1, size_t vl);
vfloat32m1_t __riscv_vfadd_tu(vfloat32m1_t vd, vfloat32m1_t vs2,
                              vfloat32m1_t vs1, unsigned int frm, size_t vl);
vfloat32m1_t __riscv_vfadd_tum(vbool32_t vm, vfloat32m1_t vd, vfloat32m1_t vs2,
                               vfloat32m1_t vs1, unsigned int frm, size_t vl);
vfloat32m1_t __riscv_vfadd_tumu(vbool32_t vm, vfloat32m1_t vd, vfloat32m1_t vs2,
                                vfloat32m1_t vs1, unsigned int frm, size_t vl);
vfloat32m1_t __riscv_vfadd_mu(vbool32_t vm, vfloat32m1_t vd, vfloat32m1_t vs2,
                              vfloat32m1_t vs1, unsigned int frm, size_t vl);
```

The naming scheme to prune everything except the instruction mnemonics is not available for all the intrinsics. Please see [Section 6.5](#) for overloaded intrinsics with irregular naming patterns.

Due to the limitations of the C language (without the aid of features like C++ templates), some intrinsics do not have an overloaded version. Therefore these intrinsics do not possess a simplified, EEW/EMUL-omitted interface. Please see [Section 6.6](#) for more detail.

6.5. Exceptions in the implicit (overloaded) naming scheme

The following intrinsics have an irregular naming pattern.

6.5.1. Widening instructions

Widening instruction intrinsics (e.g. **vwadd**) have the same return type but different types of arguments. The operand mnemonics are encoded into their overloaded versions to help distinguish them.

```
vint32m1_t __riscv_vwadd_vv(vint16mf2_t vs2, vint16mf2_t vs1, size_t vl);
vint32m1_t __riscv_vwadd_vx(vint16mf2_t vs2, int16_t rs1, size_t vl);
vint32m1_t __riscv_vwadd_wv(vint32m1_t vs2, vint16mf2_t vs1, size_t vl);
vint32m1_t __riscv_vwadd_wx(vint32m1_t vs2, int16_t rs1, size_t vl);
```

6.5.2. Type-convert instructions

Type-convert instruction intrinsics (e.g. **vfcvt.x.f**, **vfcvt.xu.f**, **vfcvt.rtz.xu.f**) encode the returning type mnemonics into their overloaded variants to help distinguish them.

The following shows how **_x**, **_rtz_x**, **_xu**, and **_rtz_xu** are appended to the suffix for distinction.

```
vint32m1_t __riscv_vfcvt_x (vfloat32m1_t src, size_t vl);
vint32m1_t __riscv_vfcvt_rtz_x (vfloat32m1_t src, size_t vl);
vuint32m1_t __riscv_vfcvt_xu (vfloat32m1_t src, size_t vl);
vuint32m1_t __riscv_vfcvt_rtz_xu (vfloat32m1_t src, size_t vl);
```

6.5.3. vreinterpret, LMUL truncate/extension, and vset/vget

These pseudo intrinsics encode the return type (e.g. **__riscv_vreinterpret_b8**) into their overloaded variants to help distinguish them.

The following shows how the return type is appended to the suffix for distinction.

```
vfloat32m1_t __riscv_vreinterpret_f32m1 (vint32m1_t src);
vuint32m1_t __riscv_vreinterpret_u32m1 (vint32m1_t src);
vint8m1_t __riscv_vreinterpret_i8m1 (vint32m1_t src);
vint16m1_t __riscv_vreinterpret_i16m1 (vint32m1_t src);
vint64m1_t __riscv_vreinterpret_i64m1 (vint32m1_t src);
vbool64_t __riscv_vreinterpret_b64 (vint32m1_t src);
vbool32_t __riscv_vreinterpret_b32 (vint32m1_t src);
vbool16_t __riscv_vreinterpret_b16 (vint32m1_t src);
vbool8_t __riscv_vreinterpret_b8 (vint32m1_t src);
vbool4_t __riscv_vreinterpret_b4 (vint32m1_t src);
```

6.6. Un-supported intrinsics for implicit (overloaded) naming scheme

Due to the limitation of the C language (without the aid of features like C++ templates), some intrinsics do not have an overloaded version. Intrinsics with characteristics of either of the following do not possess an overloaded version.

- Intrinsics with input arguments that are all scalar types and scalar types alone (e.g. unmasked vector load instruction intrinsics, `vmv.s.x`)
- Intrinsics with `vl` as the only argument (e.g. `vmc1r`, `vmset`, `vid`)
- Intrinsics with vector boolean input(s), returning a vector non-boolean vector type (e.g. `viota`)

Chapter 7. Type system

The intrinsics are designed to be strongly-typed. The intrinsics provide **vreinterpret** intrinsics to help users go across the strongly-typed scheme if necessary.

Non-mask (integer and floating-point) data types have SEW and LMUL encoded.

7.1. Integer types

Integer types have EEW and EMUL encoded into the type. The first row describes the EMUL and the first column describes the data type and element width of the scalar type.

Types with an asterisk (*) are available when **ELEN** \geq 64 (that is, unavailable under **Zve32*** and require at least **Zve64x**).

Types	EMUL=1/8	EMUL=1/4	EMUL=1/2	EMUL=1	EMUL=2	EMUL=4	EMUL=8
int8_t	vint8mf8_t*	vint8mf4_t	vint8mf2_t	vint8m1_t	vint8m2_t	vint8m4_t	vint8m8_t
int16_t	N/A	vint16mf4_t*	vint16mf2_t	vint16m1_t	vint16m2_t	vint16m4_t	vint16m8_t
int32_t	N/A	N/A	vint32mf2_t*	vint32m1_t	vint32m2_t	vint32m4_t	vint32m8_t
int64_t	N/A	N/A	N/A	vint64m1_t*	vint64m2_t*	vint64m4_t*	vint64m8_t*
uint8_t	vuint8mf8_t*	vuint8mf4_t	vuint8mf2_t	vuint8m1_t	vuint8m2_t	vuint8m4_t	vuint8m8_t
uint16_t	N/A	vuint16mf4_t*	vuint16mf2_t	vuint16m1_t	vuint16m2_t	vuint16m4_t	vuint16m8_t
uint32_t	N/A	N/A	vuint32mf2_t*	vuint32m1_t	vuint32m2_t	vuint32m4_t	vuint32m8_t
uint64_t	N/A	N/A	N/A	vuint64m1_t*	vuint64m2_t*	vuint64m4_t*	vuint64m8_t*

Table 1. Integer types

7.2. Floating-point types



*This specification uses **_Float16** to designate IEEE-754 binary16, **float** to designate IEEE-754 binary32 and **double** to designate IEEE-754 binary64.*

Floating-point types have EEW and EMUL encoded into the type. The first row describes the EMUL and the first column describes the data type and element width of the scalar type.

Floating-point types with element widths of 16 (Types=**_Float16**) require the **zvfh** and **zvfhmin** extension to be specified in the architecture.

Floating-point types with element widths of 32 (Types=**float**) require the **zve32f** extension to be specified in the architecture.

Floating-point types with element widths of 64 (Types=**double**) require the **zve64d** extension to be specified in the architecture.

Types with an asterisk (*) are available when **ELEN** \geq 64 (that is, unavailable under **Zve32f** and require at least **Zve64f**).

Types	EMUL=1/8	EMUL=1/4	EMUL=1/2	EMUL=1	EMUL=2	EMUL=4	EMUL=8
_Float16	N/A	vfloat16mf4_t*	vfloat16mf2_t	vfloat16m1_t	vfloat16m2_t	vfloat16m4_t	vfloat16m8_t
float	N/A	N/A	vfloat32mf2_t*	vfloat32m1_t	vfloat32m2_t	vfloat32m4_t	vfloat32m8_t
double	N/A	N/A	N/A	vfloat64m1_t	vfloat64m2_t	vfloat64m4_t	vfloat64m8_t

Table 2. Floating-point types

7.3. Mask types

Mask types have the ratio that is derived from **EEW/EMUL** encoded into the type. The mask types represent mask register values that follows the Mask Register Layout.

Types with an asterisk (*) are available when **ELEN** \geq 64 (that is, unavailable under **Zve32x** and require at least **Zve64x**).

Types	n = 1	n = 2	n = 4	n = 8	n = 16	n = 32	n = 64
bool	vbool1_t	vbool2_t	vbool4_t	vbool8_t	vbool16_t	vbool32_t	vbool64_t*

Table 3. Mask types

7.4. Tuple type

Tuple types encode **SEW**, **LMUL**, and **NFIELD** into the data type.

These types are utilized through the segment load/store instruction intrinsics along with getters [Section 8.5](#) and setters [Section 8.6](#) to extract/combine them. The types listed in [Section 7.1](#) and [Section 7.2](#) all have tuple types. Types under the combination of **LMUL**, **NFIELD** follows the restriction by the RVV specification, **EMUL * NFIELDS** \leq 8.

Availability of the tuple types follows the availability of their corresponding non-tuple (**NFIELD=1**) types.

Non-tuple Types (NFIELD=1)	NFIELD=2	NFIELD=3	NFIELD=4	NFIELD=5	NFIELD=6	NFIELD=7	NFIELD=8
vint8mf8_t	vint8mf8x2_t	vint8mf8x3_t	vint8mf8x4_t	vint8mf8x5_t	vint8mf8x6_t	vint8mf8x7_t	vint8mf8x8_t
vuint8mf8_t	vuint8mf8x2_t	vuint8mf8x3_t	vuint8mf8x4_t	vuint8mf8x5_t	vuint8mf8x6_t	vuint8mf8x7_t	vuint8mf8x8_t

Table 4. Tuple types (EMUL=1/8)

Non-tuple Types (NFIELD=1)	NFIELD=2	NFIELD=3	NFIELD=4	NFIELD=5	NFIELD=6	NFIELD=7	NFIELD=8
vint8mf4_t	vint8mf4x2_t	vint8mf4x3_t	vint8mf4x4_t	vint8mf4x5_t	vint8mf4x6_t	vint8mf4x7_t	vint8mf4x8_t
vuint8mf4_t	vuint8mf4x2_t	vuint8mf4x3_t	vuint8mf4x4_t	vuint8mf4x5_t	vuint8mf4x6_t	vuint8mf4x7_t	vuint8mf4x8_t
vint16mf4_t	vint16mf4x2_t	vint16mf4x3_t	vint16mf4x4_t	vint16mf4x5_t	vint16mf4x6_t	vint16mf4x7_t	vint16mf4x8_t
vuint16mf4_t	vuint16mf4x2_t	vuint16mf4x3_t	vuint16mf4x4_t	vuint16mf4x5_t	vuint16mf4x6_t	vuint16mf4x7_t	vuint16mf4x8_t
vfloat16mf4_t	vfloat16mf4x2_t	vfloat16mf4x3_t	vfloat16mf4x4_t	vfloat16mf4x5_t	vfloat16mf4x6_t	vfloat16mf4x7_t	vfloat16mf4x8_t

Table 5. Tuple types (EMUL=1/4)

Non-tuple Types (NFIELD=1)	NFIELD=2	NFIELD=3	NFIELD=4	NFIELD=5	NFIELD=6	NFIELD=7	NFIELD=8
vint8mf2_t	vint8mf2x2_t	vint8mf2x3_t	vint8mf2x4_t	vint8mf2x5_t	vint8mf2x6_t	vint8mf2x7_t	vint8mf2x8_t
vuint8mf2_t	vuint8mf2x2_t	vuint8mf2x3_t	vuint8mf2x4_t	vuint8mf2x5_t	vuint8mf2x6_t	vuint8mf2x7_t	vuint8mf2x8_t
vint16mf2_t	vint16mf2x2_t	vint16mf2x3_t	vint16mf2x4_t	vint16mf2x5_t	vint16mf2x6_t	vint16mf2x7_t	vint16mf2x8_t
vuint16mf2_t	vuint16mf2x2_t	vuint16mf2x3_t	vuint16mf2x4_t	vuint16mf2x5_t	vuint16mf2x6_t	vuint16mf2x7_t	vuint16mf2x8_t
vint32mf2_t	vint32mf2x2_t	vint32mf2x3_t	vint32mf2x4_t	vint32mf2x5_t	vint32mf2x6_t	vint32mf2x7_t	vint32mf2x8_t
vuint32mf2_t	vuint32mf2x2_t	vuint32mf2x3_t	vuint32mf2x4_t	vuint32mf2x5_t	vuint32mf2x6_t	vuint32mf2x7_t	vuint32mf2x8_t
vfloat16mf2_t	vfloat16mf2x2_t	vfloat16mf2x3_t	vfloat16mf2x4_t	vfloat16mf2x5_t	vfloat16mf2x6_t	vfloat16mf2x7_t	vfloat16mf2x8_t
vfloat32mf2_t	vfloat32mf2x2_t	vfloat32mf2x3_t	vfloat32mf2x4_t	vfloat32mf2x5_t	vfloat32mf2x6_t	vfloat32mf2x7_t	vfloat32mf2x8_t

Table 6. Tuple types (EMUL=1/2)

Non-tuple Types (NFIELD=1)	NFIELD=2	NFIELD=3	NFIELD=4	NFIELD=5	NFIELD=6	NFIELD=7	NFIELD=8
vint8m1_t	vint8m1x2_t	vint8m1x3_t	vint8m1x4_t	vint8m1x5_t	vint8m1x6_t	vint8m1x7_t	vint8m1x8_t
vuint8m1_t	vuint8m1x2_t	vuint8m1x3_t	vuint8m1x4_t	vuint8m1x5_t	vuint8m1x6_t	vuint8m1x7_t	vuint8m1x8_t
vint16m1_t	vint16m1x2_t	vint16m1x3_t	vint16m1x4_t	vint16m1x5_t	vint16m1x6_t	vint16m1x7_t	vint16m1x8_t
vuint16m1_t	vuint16m1x2_t	vuint16m1x3_t	vuint16m1x4_t	vuint16m1x5_t	vuint16m1x6_t	vuint16m1x7_t	vuint16m1x8_t
vint32m1_t	vint32m1x2_t	vint32m1x3_t	vint32m1x4_t	vint32m1x5_t	vint32m1x6_t	vint32m1x7_t	vint32m1x8_t
vuint32m1_t	vuint32m1x2_t	vuint32m1x3_t	vuint32m1x4_t	vuint32m1x5_t	vuint32m1x6_t	vuint32m1x7_t	vuint32m1x8_t
vint64m1_t	vint64m1x2_t	vint64m1x3_t	vint64m1x4_t	vint64m1x5_t	vint64m1x6_t	vint64m1x7_t	vint64m1x8_t
vuint64m1_t	vuint64m1x2_t	vuint64m1x3_t	vuint64m1x4_t	vuint64m1x5_t	vuint64m1x6_t	vuint64m1x7_t	vuint64m1x8_t
vfloat16m1_t	vfloat16m1x2_t	vfloat16m1x3_t	vfloat16m1x4_t	vfloat16m1x5_t	vfloat16m1x6_t	vfloat16m1x7_t	vfloat16m1x8_t
vfloat32m1_t	vfloat32m1x2_t	vfloat32m1x3_t	vfloat32m1x4_t	vfloat32m1x5_t	vfloat32m1x6_t	vfloat32m1x7_t	vfloat32m1x8_t
vfloat64m1_t	vfloat64m1x2_t	vfloat64m1x3_t	vfloat64m1x4_t	vfloat64m1x5_t	vfloat64m1x6_t	vfloat64m1x7_t	vfloat64m1x8_t

Table 7. Tuple types (EMUL=1)

Non-tuple Types (NFIELD=1)	NFIELD=2	NFIELD=3	NFIELD=4	NFIELD=5	NFIELD=6	NFIELD=7	NFIELD=8
vint8m2_t	vint8m2x2_t	vint8m2x3_t	vint8m2x4_t	N/A	N/A	N/A	N/A
vuint8m2_t	vuint8m2x2_t	vuint8m2x3_t	vuint8m2x4_t	N/A	N/A	N/A	N/A
vint16m2_t	vint16m2x2_t	vint16m2x3_t	vint16m2x4_t	N/A	N/A	N/A	N/A
vuint16m2_t	vuint16m2x2_t	vuint16m2x3_t	vuint16m2x4_t	N/A	N/A	N/A	N/A
vint32m2_t	vint32m2x2_t	vint32m2x3_t	vint32m2x4_t	N/A	N/A	N/A	N/A
vuint32m2_t	vuint32m2x2_t	vuint32m2x3_t	vuint32m2x4_t	N/A	N/A	N/A	N/A
vint64m2_t	vint64m2x2_t	vint64m2x3_t	vint64m2x4_t	N/A	N/A	N/A	N/A
vuint64m2_t	vuint64m2x2_t	vuint64m2x3_t	vuint64m2x4_t	N/A	N/A	N/A	N/A
vfloat16m2_t	vfloat16m2x2_t	vfloat16m2x3_t	vfloat16m2x4_t	N/A	N/A	N/A	N/A
vfloat32m2_t	vfloat32m2x2_t	vfloat32m2x3_t	vfloat32m2x4_t	N/A	N/A	N/A	N/A
vfloat64m2_t	vfloat64m2x2_t	vfloat64m2x3_t	vfloat64m2x4_t	N/A	N/A	N/A	N/A

Table 8. Tuple types (EMUL=2)

Non-tuple Types (NFIELD=1)	NFIELD=2	NFIELD=3	NFIELD=4	NFIELD=5	NFIELD=6	NFIELD=7	NFIELD=8
vint8m4_t	vint8m4x2_t	N/A	N/A	N/A	N/A	N/A	N/A
vuint8m4_t	vuint8m4x2_t	N/A	N/A	N/A	N/A	N/A	N/A
vint16m4_t	vint16m4x2_t	N/A	N/A	N/A	N/A	N/A	N/A
vuint16m4_t	vuint16m4x2_t	N/A	N/A	N/A	N/A	N/A	N/A
vint32m4_t	vint32m4x2_t	N/A	N/A	N/A	N/A	N/A	N/A
vuint32m4_t	vuint32m4x2_t	N/A	N/A	N/A	N/A	N/A	N/A
vint64m4_t	vint64m4x2_t	N/A	N/A	N/A	N/A	N/A	N/A
vuint64m4_t	vuint64m4x2_t	N/A	N/A	N/A	N/A	N/A	N/A
vfloat16m4_t	vfloat16m4x2_t	N/A	N/A	N/A	N/A	N/A	N/A
vfloat32m4_t	vfloat32m4x2_t	N/A	N/A	N/A	N/A	N/A	N/A
vfloat64m4_t	vfloat64m4x2_t	N/A	N/A	N/A	N/A	N/A	N/A

Table 9. Tuple types (EMUL=4)

Chapter 8. Pseudo intrinsics

Pseudo intrinsics provide additional utility functions to assist users in manipulating across intrinsic types. They do not map to any specific RVV instruction. The specific mapping to actual instructions is given in the description of each pseudo intrinsic.

8.1. **vsetvl**

The **vsetvl** intrinsics return the number of elements processed in a stripmining loop when provided with the element width and LMUL in the intrinsic suffix. This pseudo intrinsic is typically mapped to **vsetvli** or **vsetivli** instructions.



*The implementation must respect the ratio between SEW and LMUL given to the intrinsic. On the other hand, as mentioned in [Section 5.2](#), the **vsetvl** intrinsics do not necessarily map to the emission a **vsetvli** or **vsetivli** instruction of that exact SEW and LMUL provided. The actual value written to the **vl** control register is an implementation defined behavior and typically not known until runtime.*

8.2. **vsetvlmax**

The **vsetvlmax** intrinsics return **VLMAX** when provided with the element width and LMUL in the intrinsic suffix. This pseudo intrinsic is typically mapped to the **vsetvli** instruction.



*As mentioned in [Section 5.2](#), the **vsetvlmax** intrinsics do not necessarily map to the emission a **vsetvli** instruction of that exact SEW and LMUL provided. The actual value written to the **vl** control register is an implementation defined behavior and typically not known until runtime.*

8.3. **vreinterpret**

The **vreinterpret** intrinsics are provided for users to transition across the strongly-typed scheme. The intrinsic is limited to conversion between types operating upon the same number of registers. These intrinsics are not mapped to any instruction because reinterpretation of registers is a no-operation.

These pseudo intrinsics do not alter the bits held by a register. Please use **vfcvt/v(f)wcv** or **v(f)ncvt** intrinsics if you seek to extend, narrow, or perform real float/integer type conversions for the values.

8.4. **vundefined**

The **vundefined** intrinsics are placeholders to represent unspecified values in variable initialization, or as arguments of **vset** and **vcreate**. These pseudo intrinsics are not mapped to any instruction.

8.5. **vget**

The **vget** intrinsics allow users to obtain small LMUL values from larger LMUL ones. The **vget** intrinsics also allows users to extract non-tuple (**NFIELD=1**) types from tuple (**NFIELD>1**) types after segment load intrinsics. The index provided must be a constant known at compile time.

These pseudo intrinsics do not map to any real instruction. The compiler may emit zero or more instructions to implement the semantics of these pseudo intrinsics. The precise set of instructions emitted is a compiler optimization issue.

8.6. **vset**

The **vset** intrinsics allow users to combine small LMUL values into larger LMUL ones. The **vset** intrinsics also allows users to combine non-tuple (**NFIELD=1**) types to tuple (**NFIELD>1**) types for segment store intrinsics. The index provided must be a constant known at compile time.

These pseudo intrinsics do not map to any real instruction. The compiler may emit zero or more instructions to implement the semantics of these pseudo intrinsics. The precise set of instructions emitted is a compiler optimization issue.

8.7. **vlmul_trunc**

The **vlmul_trunc** intrinsics are syntactic sugar for RVV vector types other than tuples and have the same semantic as **vget** with the **index** operand having the value **0**.

8.8. **vlmul_ext**

The **vlmul_ext** intrinsics are syntactic sugar for RVV vector types other than tuples and have the same semantic as **vset** with the **index** operand having the value **0**.

8.9. **vcreate**

The **vcreate** intrinsics are syntactic sugar for the creation of values of RVV types. They have the same semantic as multiple **vset** pseudo intrinsics filling in values accordingly.

*Example 1. Pseudo intrinsic **vcreate** used to build a SEW=32, LMUL=4 **float** vector (**vfloat32m4**) from two SEW=32, LMUL=2 **float** vectors (**vfloat32m2**).*

```
// Given the following declarations
vfloat32m4_t dest;
vfloat32m2_t v0 = ...;
vfloat32m2_t v1 = ...;

// this pseudo intrinsic
dest = __riscv_vcreate_v_f32m2_f32m4(v0, v1);

// is semantically equivalent to
dest = __riscv_vset_v_f32m2_f32m4(__riscv_vundefined_f32m4(), 0, v0);
dest = __riscv_vset_v_f32m2_f32m4(dest, 1, v1);
```

8.10. **vlenb**

The **vlenb** intrinsic returns what is held inside the read-only CSR **vlenb**, which is the vector register length in bytes. This pseudo intrinsic is mapped to a **csrr** instruction that reads from the CSR **vlenb**.

```
unsigned __riscv_vlenb();
```

Chapter 9. Programming Notes

9.1. Agnostic value in the RVV C intrinsics

An agnostic value is in the granularity of "element" contained in an RVV type. Agnostic values are either:

- Tail element(s) of an RVV value produced through a **ta** instruction
- Masked-off element(s) of an RVV value produced through a **ma** instruction
- All elements in an uninitialized value
- A value assigned with the **vundefined** intrinsics

An agnostic value is an indeterminate value and evaluation of an agnostic value is undefined behavior. Users should not rely on any evaluation to an agnostic value.

9.2. Copying vector register group contents

There is no intrinsic that directly maps to the whole vector register move instructions (**vmv<nr>r.v**).

For copying of the vector contents in whole, we encourage the users to use the assignment operator (**=**).

The assignment operator (**=**) represents the semantic of a whole vector register (group) copy for the expression on the right hand side to the RVV type object on the left hand side. The semantic will still maintain a whole vector register content copy for fractional LMUL types. This enables the compiler to coalesce register usage when possible.

Users may leverage the vector move intrinsics (**vmv_v_v**) intrinsics if they hope to copy vector register groups with **vl != VLMAX**.

9.3. The passthrough (**vd**) argument in the intrinsics

Intrinsics whose computation is relevant to the value held in destination register (assembly mnemonics **vd**) have a **vd** argument in them. The following list enumerates the intrinsics that have a **vd** argument. Please see the appendix for the exact prototypes of these intrinsics.

- Intrinsics with tail-undisturbed (**vta=0**)
- Intrinsics with mask-undisturbed (**vma=0**)
- Intrinsics representing Vector Multiply-Add Operations
- Intrinsics representing Vector Slideup Instructions

For intrinsics with no **vd** argument, the implementation is free to pick any register as the destination register.

9.4. Assumption of **vstart=0** for intrinsics users.

The **vstart** CSR is currently not exposed to the intrinsics programmer, and the intrinsics have the semantics of **vstart = 0**. Support for positive **vstart** values is implementation -defined; thus,

portable application software should not set `vstart > 0`.

9.5. Assembly generated from the intrinsics

Some users may expect the intrinsics to directly translate and appear in the assembly; however, the intrinsics are the interfaces that expose the vector instruction semantics. The implementation is free to optimize them out if there is an opportunity.

9.6. Bookkeeping of configurations

Control of `vl`, `vtype`, `vxrm`, and `frm` is not directly exposed to the user. The implementation is responsible for setting the correct values into these CSRs to achieve the expected semantics of the intrinsic functions with respect to the conventions defined in the ISA specification ([RISC-V "V" Vector Extension, n.d.](#)) and ABI specification ([RISC-V Calling Conventions: Vector, n.d.](#)).

9.7. Strided load/store with stride of 0

The RVV specification mentions that the strided load/store instruction with a stride of 0 could have different behaviors, performing all memory accesses or fewer memory operations. Since needing all memory accesses isn't likely to be common, the implementation is allowed to generate fewer memory operations with strided load/store intrinsics.

In other words, the compiler does not guarantee generating the instruction for all memory accesses in strided load/store intrinsics with a stride of 0. If the user needs all memory accesses to be performed, they should use an indexed load/store intrinsics with all zero indices.

9.8. Leveraging instructions with operand mnemonics of `vi`

The intrinsics provide variants with operand mnemonics of `vv` and `vx`, but not `vi`. This was an intentional design to reduce the total amount of out-going intrinsics.

It is an optimization issue for the implementation to emit instructions with operand mnemonics of `vi` when an immediate that can be expressed within 5-bit is provided to the intrinsics.

9.9. Mixing inline assembly and intrinsics

The compiler will be conservative to registers (`vtype`, `vxrm`, `frm`) when encountering inline assembly. Users should be aware that mixing uses of intrinsics and inline assembly will result in extra save and restore.

9.10. The `new_vl` argument in fault-only-first load intrinsics

The fault-only-first load intrinsics write the new value inside the `vl` register into the address of the `new_vl` argument. Providing an illegal memory location is undefined behavior. == Intrinsics for BFloat16 (Brain Float 16) instruction set extensions

The RISC-V vector C intrinsics supports intrinsics that exposes the control of BFloat16 (Brain Float 16) instruction set extensions³¹.

9.11. Naming scheme

The BFloat16 intrinsics follows the naming scheme defined under [Chapter 6](#), with **bf** as the abbreviation for BFloat16 types in the function suffix.

9.12. Control of the vector extension programming model

The BFloat16 intrinsics follows provides the same control of the vector programming model defined under [Chapter 5](#). Intrinsics that represents BFloat16 instructions that are affected by **frm** (**vfncvtfbf16.f.f.w** and **vfwmacbf16**) follow what is defined under [Section 5.6](#) and provides variants of [\[implicit-frm\]](#) and [Section 5.6.2](#).

9.13. Type system

Floating-point types have EEW and EMUL encoded into the type. The first row describes the EMUL and the first column describes the data type and element width of the scalar type.

Floating-point types with element widths of 16 (Types=**__bf16**) require the **zfbfmin** and **zvfbfmin** extension to be specified in the architecture.



Although C++23 introduces `<stdfloat>` for fixed-width floating-point types, this latest standard is not yet supported in the upstream RISC-V compiler. The specification (along with the prototype lists in appendix) uses **__bf16** to represent the BFloat16 floating-point type.

Types	EMUL=1/8	EMUL=1/4	EMUL=1/ 2	EMUL=1	EMUL=2	EMUL=4	EMUL=8
__bf16	N/A	<code>vbfloat16mf4_t</code>	<code>vbfloat16mf2_t</code>	<code>vbfloat16m1_t</code>	<code>vbfloat16m2_t</code>	<code>vbfloat16m4_t</code>	<code>vbfloat16m8_t</code>

Table 10. BFloat16 types

Non-tuple Types (NFIELD=1)	NFIELD=2	NFIELD=3	NFIELD=4	NFIELD=5	NFIELD=6	NFIELD=7	NFIELD=8
<code>vbfloat16mf4_t</code>	<code>vbfloat16mf4x2_t</code>	<code>vbfloat16mf4x3_t</code>	<code>vbfloat16mf4x4_t</code>	<code>vbfloat16mf4x5_t</code>	<code>vbfloat16mf4x6_t</code>	<code>vbfloat16mf4x7_t</code>	<code>vbfloat16mf4x8_t</code>
<code>vbfloat16mf2_t</code>	<code>vbfloat16mf2x2_t</code>	<code>vbfloat16mf2x3_t</code>	<code>vbfloat16mf2x4_t</code>	<code>vbfloat16mf2x5_t</code>	<code>vbfloat16mf2x6_t</code>	<code>vbfloat16mf2x7_t</code>	<code>vbfloat16mf2x8_t</code>
<code>vbfloat16m1_t</code>	<code>vbfloat16m1x2_t</code>	<code>vbfloat16m1x3_t</code>	<code>vbfloat16m1x4_t</code>	<code>vbfloat16m1x5_t</code>	<code>vbfloat16m1x6_t</code>	<code>vbfloat16m1x7_t</code>	<code>vbfloat16m1x8_t</code>
<code>vbfloat16m2_t</code>	<code>vbfloat16m2x2_t</code>	<code>vbfloat16m2x3_t</code>	<code>vbfloat16m2x4_t</code>	N/A	N/A	N/A	N/A
<code>vbfloat16m4_t</code>	<code>vbfloat16m4x2_t</code>	N/A	N/A	N/A	N/A	N/A	N/A

Table 11. Tuple types

9.14. Psuedo intrinsics

The RISC-V vector BFloat16 types (provided under [Section 9.13](#)) also have pseudo intrinsics variants from [Chapter 8](#) to help variable declaration and manipulation across intrinsic types.

Chapter 10. References

⁰[Github - riscv/riscv-v-spec/v-spec.adoc](#)



*Standard extensions are merged into **riscv/riscv-isa-manual** after ratification. There is an on-going pull request ²⁶ for the "V" extension to be merged. At this moment this intrinsics specification still references the frozen draft ⁰. This reference will be updated in the future once the pull request has been merged.*

¹[Github - riscv-non-isa/riscv-c-api-doc/riscv-c-api.md](#)

²[User Guide for RISC-V Target](#)

³[RISC-V Options \(Using the GNU Compiler Collection \(GCC\)\)](#)

⁴Section 3.4.1 (Vector selected element width **vsew[2:0]**) in the specification ⁰

⁵Section 3.4.2 (Vector Register Grouping (**vlmul[2:0]`**)) in the specification ⁰

⁶Section 3.4.3 (Vector Tail Agnostic and Vector Mask Agnostic **vta** and **vma**) in the specification ⁰

⁷Section 5.3 (Vector Masking) in the specification ⁰

⁸Section 3.8 (Vector Fixed-Point Rounding Mode Register **vxrm**) in the specification ⁰

⁹[psABI: Vector Register Convention](#)

¹⁰[The RISC-V Instruction Set Manual: 8.2 Floating-Point Control and Status Register](#)

¹¹Section 3.5 (Vector Length Register) in the specification ⁰

¹²Section 3.4.2 in the specification ⁰

¹³Section 11.13, 11.14, 13.6, 13.7 in the specification ⁰

¹⁴Section 4.5 (Mask Register Layout) in the specification ⁰

¹⁵Section 7.5 in the specification ⁰

¹⁶Section 7.8 in the specification ⁰

¹⁷Section 5.2 (Vector Operands) in the specification ⁰

¹⁸Section 6 (Configuration-Setting Instructions) in the specification ⁰

¹⁹Section 18 (Standard Vector Extensions) in the specification ⁰

²⁰Section 18.2 (Zve*: Vector Extensions for Embedded Processors) in the specification ⁰

²¹Section 12 (Vector Fixed-Point Arithmetic Instructions) in the specification ⁰

²²Section 3.9 (3.9. Vector Fixed-Point Saturation Flag **vxsat**) in the specification ⁰

²³Section 13 (Vector Floating-Point Instructions) in the specification ⁰

²⁴Section 16.3.1 (Vector Slideup Instructions) in the specification ⁰

²⁵Section 3.7 (Vector Start Index CSR **vstart**) in the specification ⁰

²⁶[riscv/riscv-isa-manual#1088](#)

²⁷Section 6.3 (Constraints on Setting **vl**) in the specficiation ⁰

²⁸Section 6.4 (Example of stripmining and changes to SEW) in the specification ⁰

²⁹Section 3.6 (Vector Byte Length **vlenb**) in the specification ⁰

³⁰Section 16.6 (Whole Vector Register Move) in the specification ⁰

³¹[RISC-V BFloat16 Specification](#)

Chapter 11. Examples



This section is non-normative.



No claims about efficiency are made the examples presented in this section.

This section presents examples that use the RVV intrinsics specified in this document. The examples are in C and assume `#include <riscv_vector.h>` has appeared earlier in the source code.

11.1. Memory copy

*Example 2. An implementation of the **memcpy** function of the C Standard library using RVV intrinsics.*

```
void *memcpy_rvv(void *restrict destination, const void *restrict source,
    size_t n) {
    unsigned char *dst = destination;
    const unsigned char *src = source;
    // copy data byte by byte
    for (size_t vl; n > 0; n -= vl, src += vl, dst += vl) {
        vl = __riscv_vsetvl_e8m8(n);
        // Load src[0..vl)
        vuint8m8_t vec_src = __riscv_vle8_v_u8m8(src, vl);
        // Store dst[0..vl)
        __riscv_vse8_v_u8m8(dst, vec_src, vl);
        // src is incremented vl (bytes)
        // dst is incremented vl (bytes)
        // n is decremented vl
    }
    return destination;
}
```

11.2. SAXPY

Consider the following function that implements a SAXPY-like kernel.

```
void saxpy_reference(size_t n, const float a, const float *x, float *y) {
    for (size_t i = 0; i < n; ++i) {
        y[i] = a * x[i] + y[i];
    }
}
```

Example 3. An implementation of SAXPY using RVV intrinsics.

```
void saxpy_rvv(size_t n, const float a, const float *x, float *y) {
    for (size_t vl; n > 0; n -= vl, x += vl, y += vl) {
        vl = __riscv_vsetvl_e32m8(n);
        // Load x[i..i+vl)
        vfloat32m8_t vx = __riscv_vle32_v_f32m8(x, vl);
        // Load y[i..i+vl)
        vfloat32m8_t vy = __riscv_vle32_v_f32m8(y, vl);
        // Computes vy[0..vl) + a*vx[0..vl)
        // and stores it in y[i..i+vl)
        __riscv_vse32_v_f32m8(y, __riscv_vfmacc_vf_f32m8(vy, a, vx, vl), vl);
    }
}
```

}

11.3. Matrix multiplication

Consider the following function that implements a naive matrix multiplication.

```
// matrix multiplication
// C[0..n][0..m] = A[0..n][0..p] x B[0..p][0..m]
void matmul_reference(double *a, double *b, double *c, int n, int m, int p) {
    for (int i = 0; i < n; ++i)
        for (int j = 0; j < m; ++j) {
            c[i * m + j] = 0;
            for (int k = 0; k < p; ++k) {
                c[i * m + j] += a[i * p + k] * b[k * m + j];
            }
        }
}
```

The following example is a version of the matrix multiplication. The accumulation on `c[i * m + j]` is implemented using partial accumulations followed by a single final accumulation.

Example 4. An implementation of a naive matrix multiplication using RVV intrinsics.

```
void matmul_rvv(double *a, double *b, double *c, int n, int m, int p) {
    size_t vlmax = __riscv_vsetvln_e64m1();
    for (int i = 0; i < n; ++i)
        for (int j = 0; j < m; ++j) {
            double *ptr_a = &a[i * p];
            double *ptr_b = &b[j];
            int k = p;
            // Set accumulator to zero.
            vfloat64m1_t vec_s = __riscv_vfmv_v_f_f64m1(0.0, vlmax);
            vfloat64m1_t vec_zero = __riscv_vfmv_v_f_f64m1(0.0, vlmax);
            for (size_t vl; k > 0; k -= vl, ptr_a += vl, ptr_b += vl * m) {
                vl = __riscv_vsetvl_e64m1(k);

                // Load row a[i][k..k+vl)
                vfloat64m1_t vec_a = __riscv_vle64_v_f64m1(ptr_a, vl);
                // Load column b[k..k+vl)[j]
                vfloat64m1_t vec_b =
                    __riscv_vlse64_v_f64m1(ptr_b, sizeof(double) * m, vl);

                // Accumulate dot product of row and column. If vl < vlmax we need to
                // preserve the existing values of vec_s, hence the tu policy.
                vec_s = __riscv_vfmacc_vv_f64m1_tu(vec_s, vec_a, vec_b, vl);
            }

            // Final accumulation.
            vfloat64m1_t vec_sum =
                __riscv_vfredusum_vs_f64m1_f64m1(vec_s, vec_zero, vlmax);
            double sum = __riscv_vfmv_f_s_f64m1_f64(vec_sum);
            c[i * m + j] = sum;
        }
}
```

11.4. String copy

Example 5. An implementation of the **strcpy** function of the C Standard Library using RVV intrinsics.

```
char *strcpy_rvv(char *destination, const char *source) {
    unsigned char *dst = (unsigned char *)destination;
    unsigned char *src = (unsigned char *)source;
    size_t vlmax = __riscv_vsetvmax_e8m8();
    long first_set_bit = -1;

    // This loop stops when among the loaded bytes we find the null byte
    // of the string i.e., when first_set_bit >= 0
    for (size_t vl; first_set_bit < 0; src += vl, dst += vl) {
        // Load up to vlmax elements if possible.
        // vl is set to the maximum number of elements, no more than vlmax, that
        // could be loaded without causing a memory fault.
        vuint8m8_t vec_src = __riscv_vle8ff_v_u8m8(src, &vl, vlmax);

        // Mask that states where null bytes are in the loaded bytes.
        vbool1_t string_terminate = __riscv_vmseq_vx_u8m8_b1(vec_src, 0, vl);

        // If the null byte is not in the loaded bytes the resulting mask will
        // be all ones, otherwise only the elements up to and including the
        // first null byte of the resulting will be enabled.
        vbool1_t mask = __riscv_vmsif_m_b1(string_terminate, vl);

        // Store the enabled elements as determined by the mask above.
        __riscv_vse8_v_u8m8_m(mask, dst, vec_src, vl);

        // Determine if we found the null byte in the loaded bytes.
        // If not found, first_set_bit is set to all ones (i.e., -1), otherwise
        // first_set_bit will be the number of the first element enabled in the
        // mask.
        first_set_bit = __riscv_vfirst_m_b1(string_terminate, vl);
    }
    return destination;
}
```

11.5. Control flow

Consider the following function that computes the division of two arrays elementwise but sets the result to a given value when the element of the divisor array is zero.

```
void branch_ref(double *a, double *b, double *c, int n, double constant) {
    for (int i = 0; i < n; ++i) {
        c[i] = (b[i] != 0.0) ? a[i] / b[i] : constant;
    }
}
```

The following example applies if-conversion using masks to implement the semantics of the conditional operator.

Example 6. An implementation of **branch_ref** using RVV intrinsics.

```
void branch_rvv(double *a, double *b, double *c, int n, double constant) {
    // set vlmax and initialize variables
    size_t vlmax = __riscv_vsetvmax_e64m1();
```

```

// "Broadcast" the value of constant to all (vlmax) the elements in
// vec_constant.
vfloat64m1_t vec_constant = __riscv_vfmv_v_f64m1(constant, vlmax);
for (size_t vl; n > 0; n -= vl, a += vl, b += vl, c += vl) {
    vl = __riscv_vsetvl_e64m1(n);

    // Load a[i..i+vl)
    vfloat64m1_t vec_a = __riscv_vle64_v_f64m1(a, vl);
    // Load b[i..i+vl)
    vfloat64m1_t vec_b = __riscv_vle64_v_f64m1(b, vl);

    // Compute a mask whose enabled elements will correspond to the
    // elements of b that are not zero.
    vbool64_t mask = __riscv_vmfne_vf_f64m1_b64(vec_b, 0.0, vl);

    // Use mask undisturbed policy to compute the division for the
    // elements enabled in the mask, otherwise set them to the given
    // constant above (maskedoff).
    vfloat64m1_t vec_c = __riscv_vfdiv_vv_f64m1_mu(
        mask, /*maskedoff*/ vec_constant, vec_a, vec_b, vl);

    // Store into c[i..i+vl)
    __riscv_vse64_v_f64m1(c, vec_c, vl);
}
}

```

11.6. Reduction and counting

Consider the following function that computes the dot product of two arrays excluding elements of the first array (along with the correspondign element of the second array) where the value is 42. The function also counts how many pairs of elements took part in the dot-product.

```

void reduce_reference(double *a, double *b, double *result_sum,
                    int *result_count, int n) {
    int count = 0;
    double s = 0.0;
    for (int i = 0; i < n; ++i) {
        if (a[i] != 42.0) {
            s += a[i] * b[i];
            count++;
        }
    }

    *result_sum = s;
    *result_count = count;
}

```

The following example implements the accumulation of the `s` variable doing several partial accumulations followed by a final accumulation.

Example 7. An implementation of `reduce_reference` using RVV intrinsics.

```

void reduce_rvv(double *a, double *b, double *result_sum, int *result_count,
               int n) {
    int count = 0;
    // set vlmax and initialize variables
    size_t vlmax = __riscv_vsetvlmax_e64m1();
    vfloat64m1_t vec_zero = __riscv_vfmv_v_f64m1(0.0, vlmax);
    vfloat64m1_t vec_s = __riscv_vfmv_v_f64m1(0.0, vlmax);

```

```

for (size_t vl; n > 0; n -= vl, a += vl, b += vl) {
    vl = __riscv_vsetvl_e64m1(n);

    // Load a[i..i+vl)
    vfloat64m1_t vec_a = __riscv_vle64_v_f64m1(a, vl);
    // Load b[i..i+vl)
    vfloat64m1_t vec_b = __riscv_vle64_v_f64m1(b, vl);

    // Compute a mask whose enabled elements will correspond to the
    // elements of a that are not 42.
    vbool64_t mask = __riscv_vmfne_vf_f64m1_b64(vec_a, 42.0, vl);

    // for all e in [0..vl)
    //  vec_s[e] ← vec_s[e] + vec_a[e] * vec_b[e], if mask[e] is enabled
    //              vec_s[e]                , otherwise (mask undisturbed)
    vec_s = __riscv_vfmacc_vv_f64m1_tumu(mask, vec_s, vec_a, vec_b, vl);

    // Adds to count the number of elements in mask that are enabled.
    count += __riscv_vcpop_m_b64(mask, vl);
}

vfloat64m1_t vec_sum;
// Final accumulation.
vec_sum = __riscv_vfredusum_vs_f64m1_f64m1(vec_s, vec_zero, vlmax);
double sum = __riscv_vfmv_f_s_f64m1_f64(vec_sum);

// Return values.
*result_sum = sum;
*result_count = count;
}

```